

Week 5

Course. Introduction to Machine Learning **Theory 5. Visualization**

Dr. Maria Salamó Llorente
maria.salamo@ub.edu

Dept. Mathematics and Informatics,
Faculty of Mathematics and Informatics,
University of Barcelona (UB)

Introduction to Machine Learning

Unsupervised
Learning

Supervised
Learning

Decision
Learning
Theory

Cluster
Analysis

Factor
Analysis

Visuali-
zation

Non Linear Decision

Linear
Decision

Basic
concepts
of
Decision
Learning
Theory

K-Means ,
EM

PCA, ICA

SOM,
MDS

Lazy
Learning
(K-NN, IBL,
CBR)

Overfitting,
model
selection and
feature
selection

Kernel
Learning

Ensemble
Learning
(NN, Trees,
Adaboost)

Perceptron,
SVM

Bias/Variance,
VC dimension,
Practical
advice of how
to use learning
algorithms

Contents

1. Introduction to visualization

2. Self-Organizing Maps

1. Introduction to Self-Organizing Maps

2. Defining SOM : Output layers, topological structure, neurons, the adaptation in SOM, training process,

3. SOM algorithm

4. Example

5. Summary

3. Multidimensional Scaling

1. Introduction to multidimensional scaling

2. Taxonomy of MDS

3. Classical MDS algorithm

4. Example

5. Summary

Introduction to Visualization

- **Objective:** Visualizing low-dimensional views of high-dimensional data
- This lecture is mainly devoted to:
 - Self-Organizing Maps
 - Multi-Dimensional Scaling

Keywords

Consider the following ideas

knowledge

data

wisdom

information

Exercise

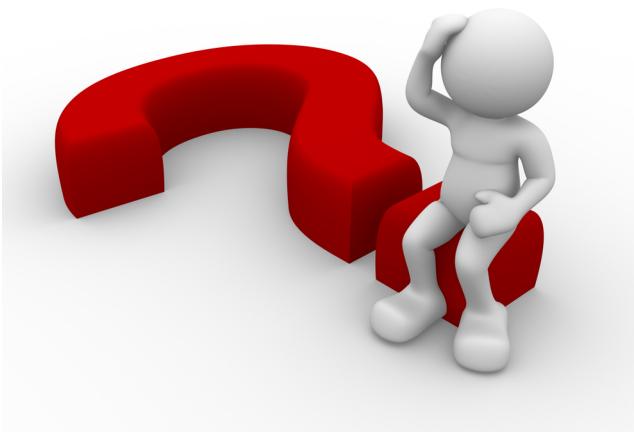
- Now, order previous terms as logically as possible

_____ > _____ > _____ > _____

Take a moment to think about it

....

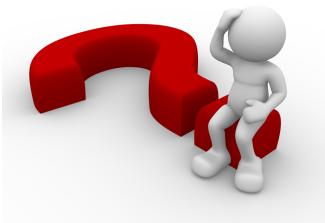
make a note that explain your decision
(You will have the answer in the online session)





Connect keywords

*Connect each keyword
with its corresponding image*



data



information



knowledge



wisdom



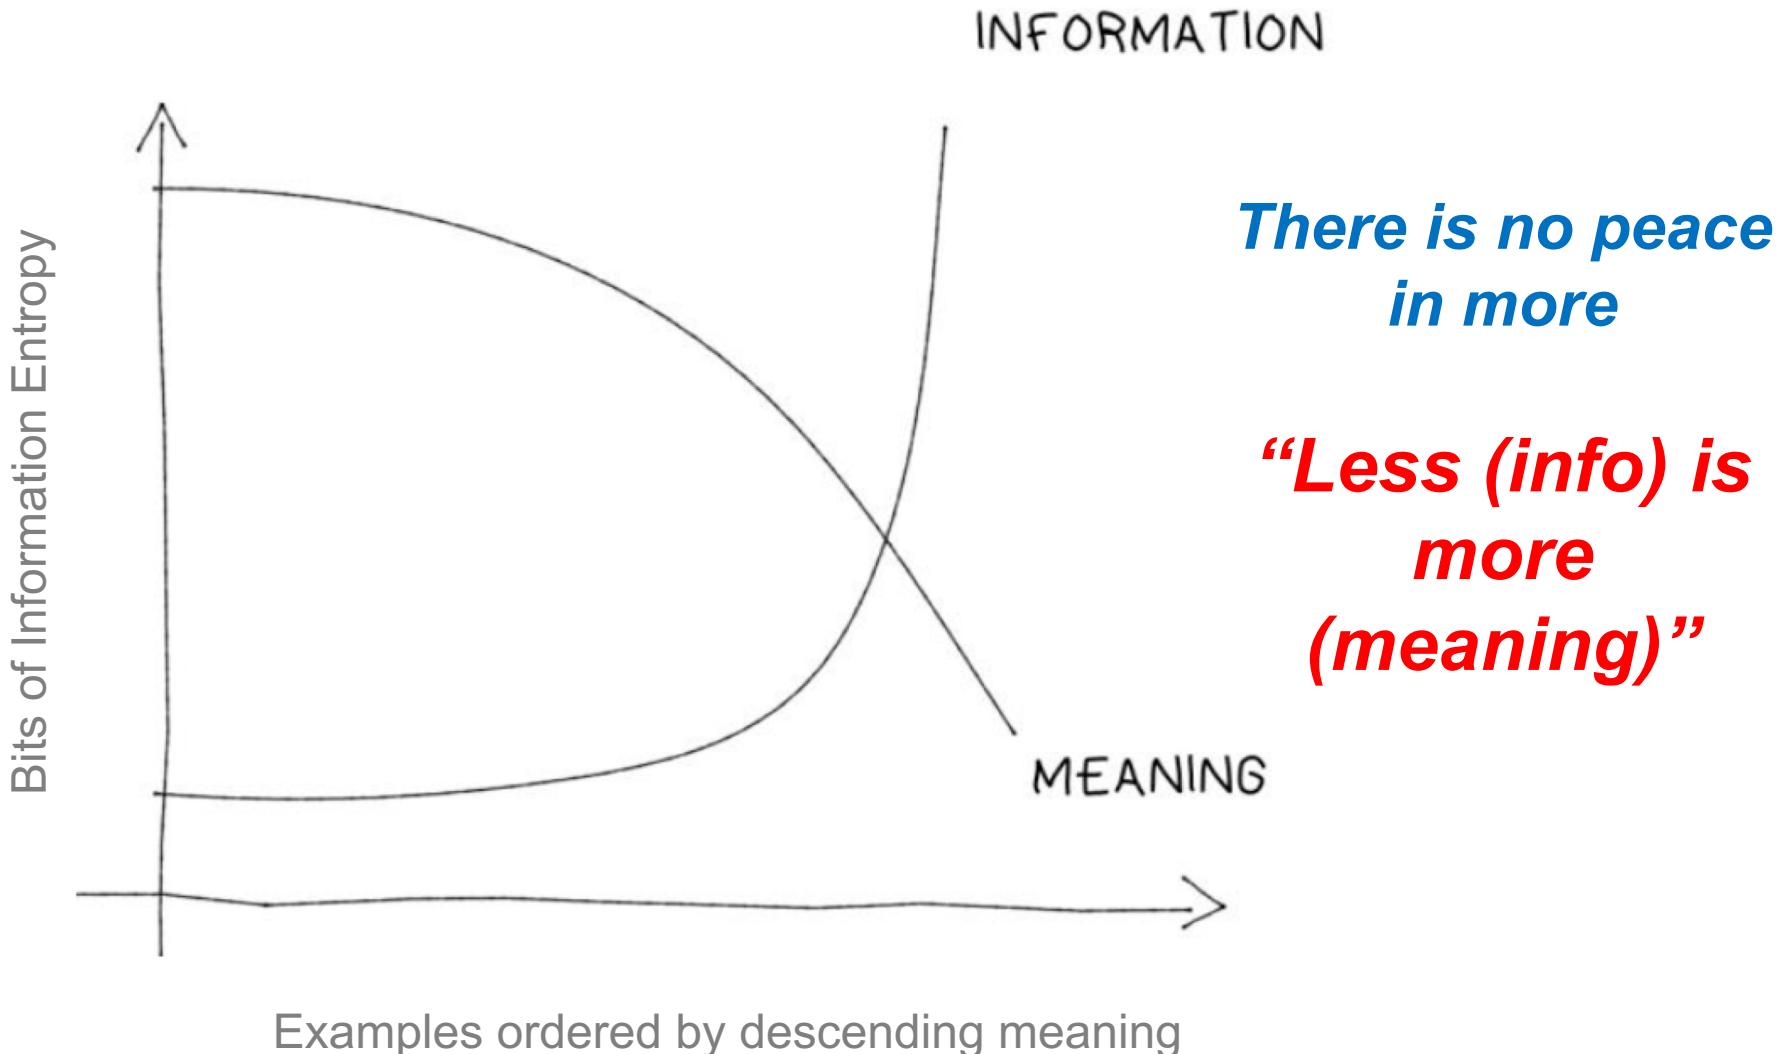


How is wisdom made?

Data → Wisdom
How?

*One way to arrive at wisdom is the **Synthesis** process*

Information vs Meaning

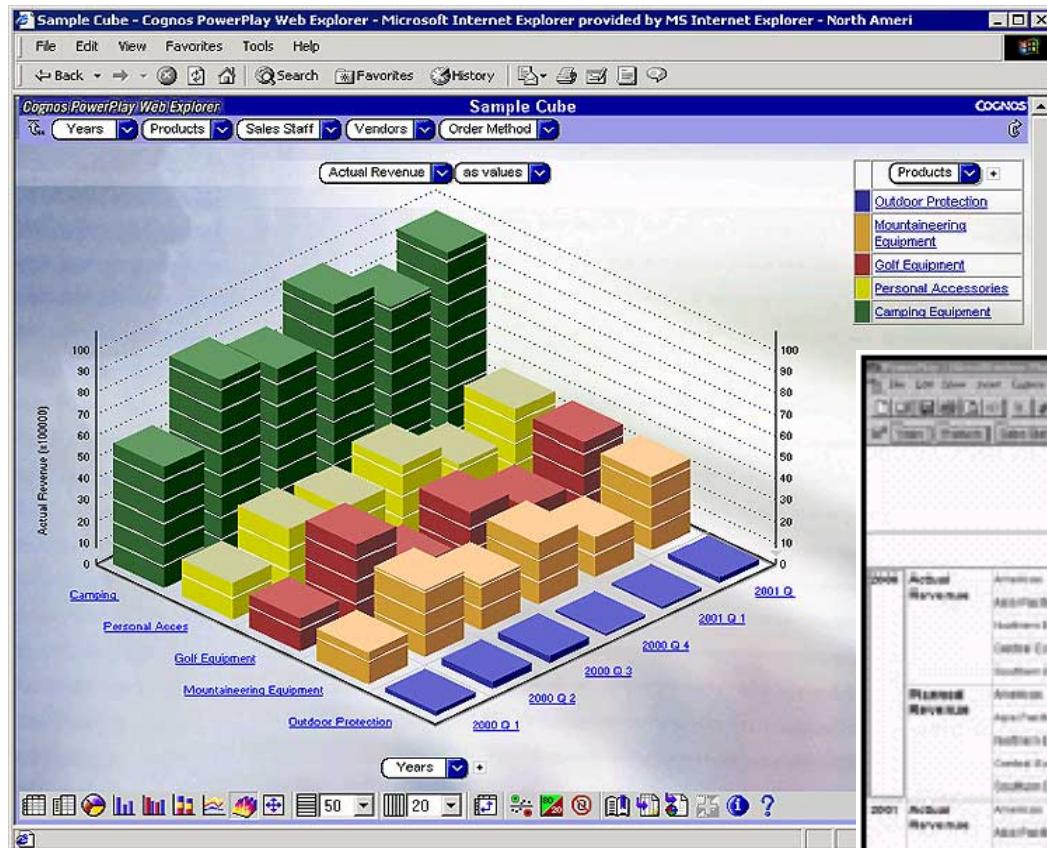
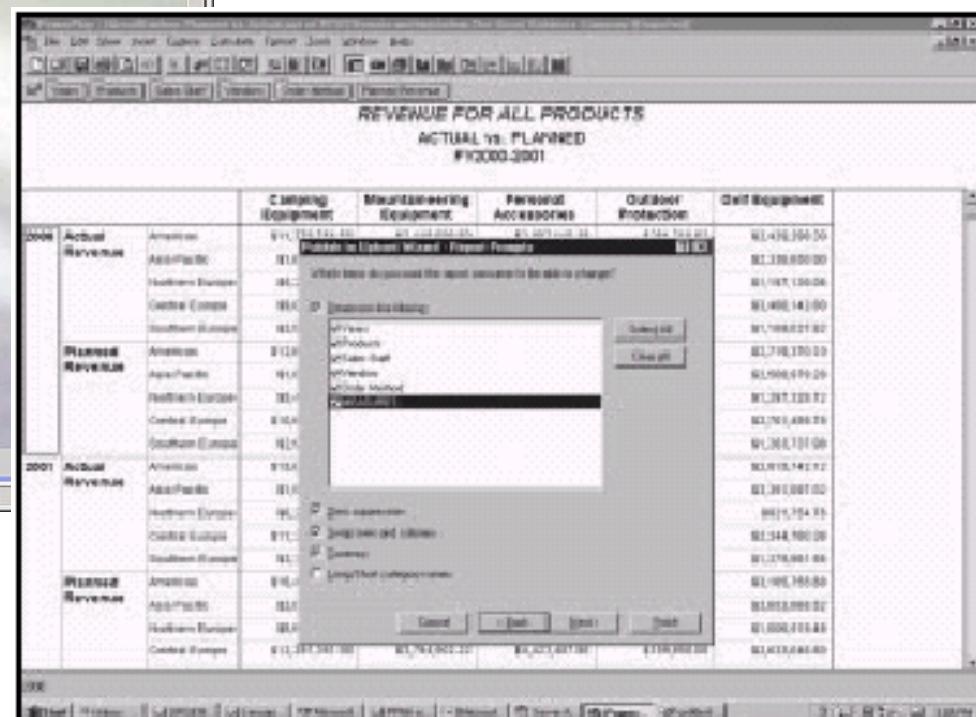


Visualization

VISUALIZATION is the study of the visual representation of the data

Main goal is to communicate the information clearly and effectively through graphical means

Data visualization vs model visualization

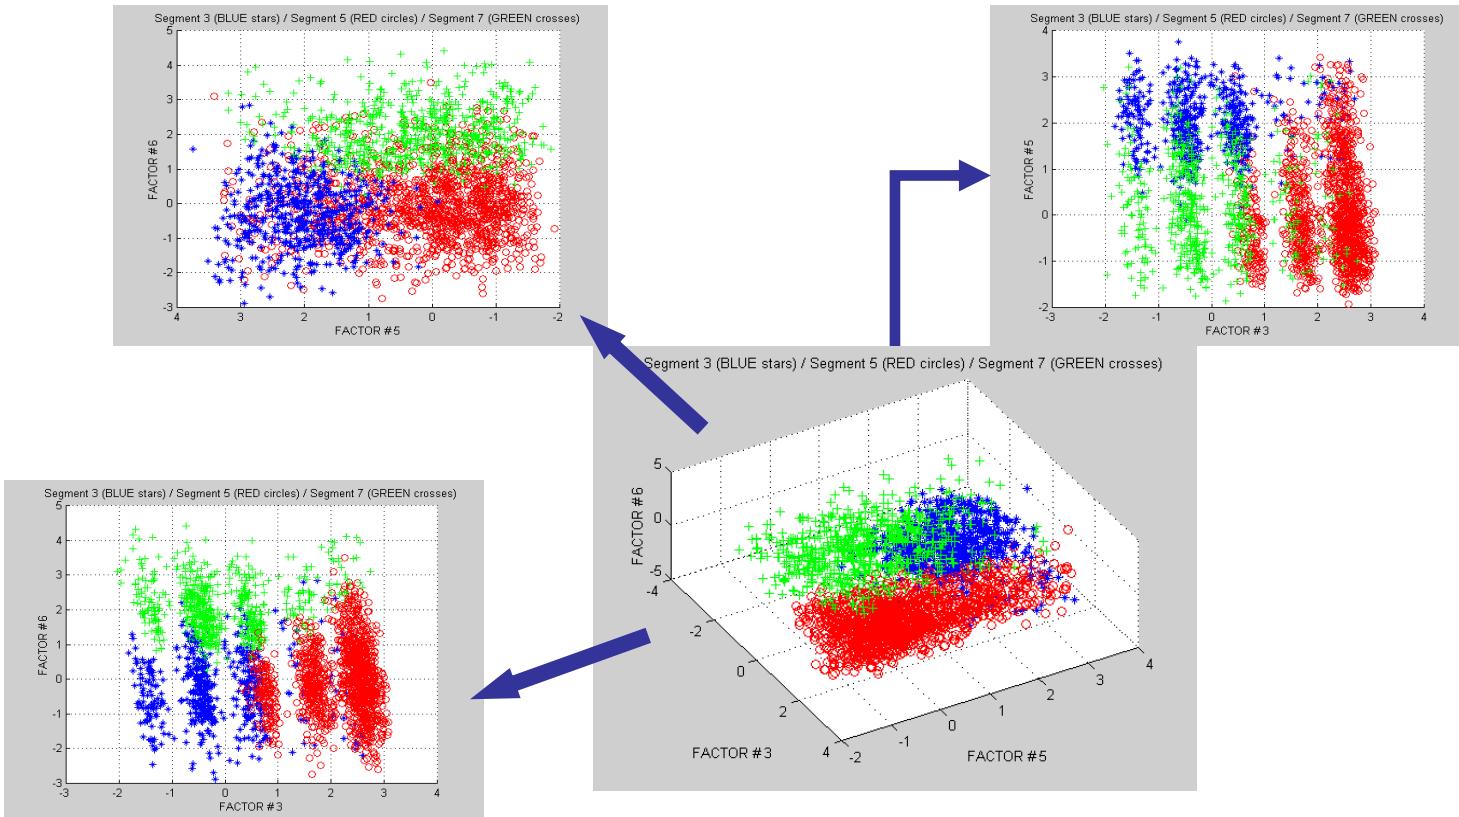
The figure shows a spreadsheet titled "REVENUE FOR ALL PRODUCTS" for the period "ACTUAL vs. PLANNED #10000-2001". The table has columns for Product Category, Region, and Revenue Type. The rows show data for 2000 and 2001, broken down by quarter and specific product types like "Camping Equipment", "Mountaineering Equipment", "Personal Accessories", "Outdoor Protection", and "Golf Equipment". The table includes various formulas and dropdown menus for data entry.

	Camping Equipment	Mountaineering Equipment	Personal Accessories	Outdoor Protection	Golf Equipment
2000 Actual Revenue	\$100,000	\$80,000	\$60,000	\$40,000	\$20,000
2000 Planned Revenue	\$120,000	\$90,000	\$70,000	\$50,000	\$30,000
2001 Actual Revenue	\$110,000	\$95,000	\$75,000	\$55,000	\$35,000
2001 Planned Revenue	\$130,000	\$100,000	\$80,000	\$60,000	\$40,000

We can visualize the data using a tool,
here an OLAP tool

Data visualization vs model visualization

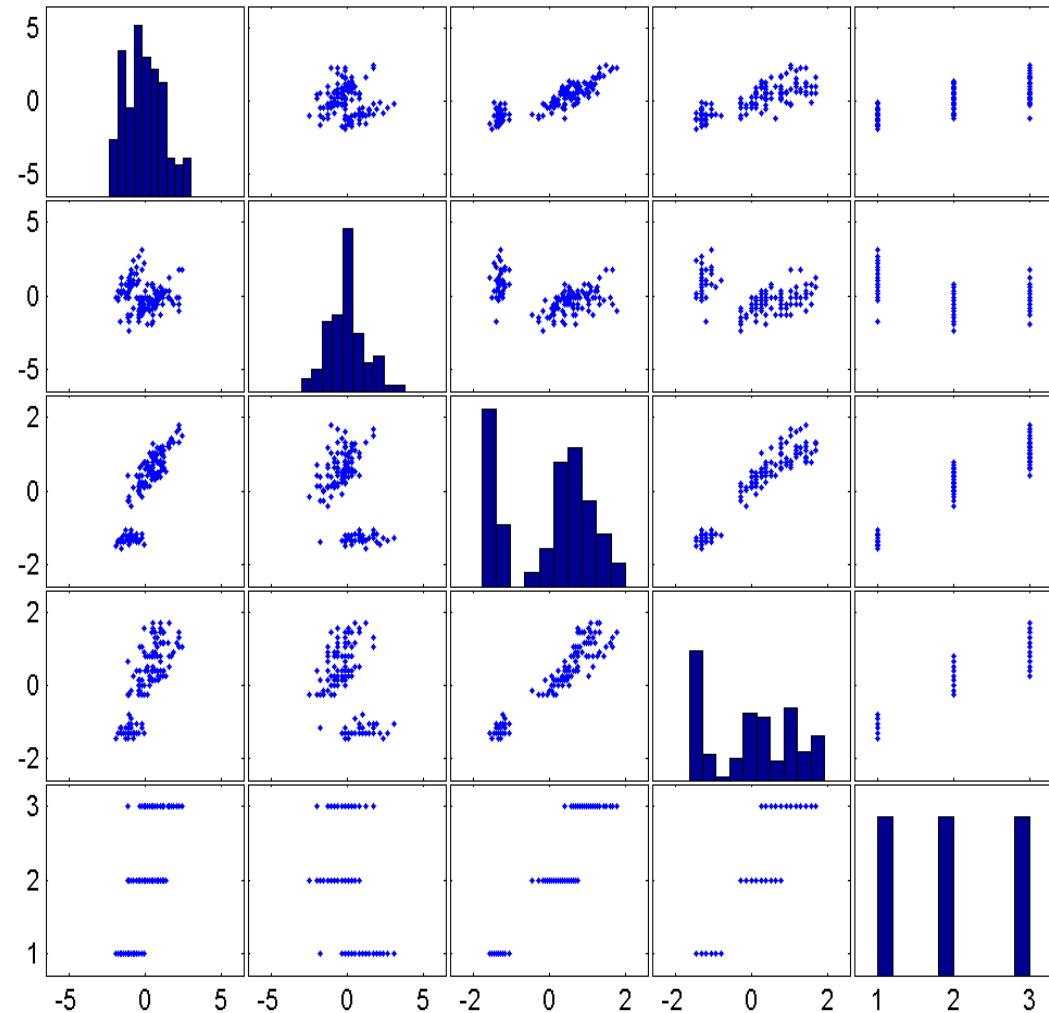
*... or we can visualize a **model**. Here we visualize in 3D three factors of interest in a factor analysis solution*



How do we visualize data of high (or even very high) dimensionality?

- **Eliminate dimensions** (data variables): those which are redundant and/or uninformative (at least you manage to alleviate part of the problem...)
- **Divide & conquer**: a classic approach where you create multiple visualizations of low dimensionality
- **Latent and projection models**

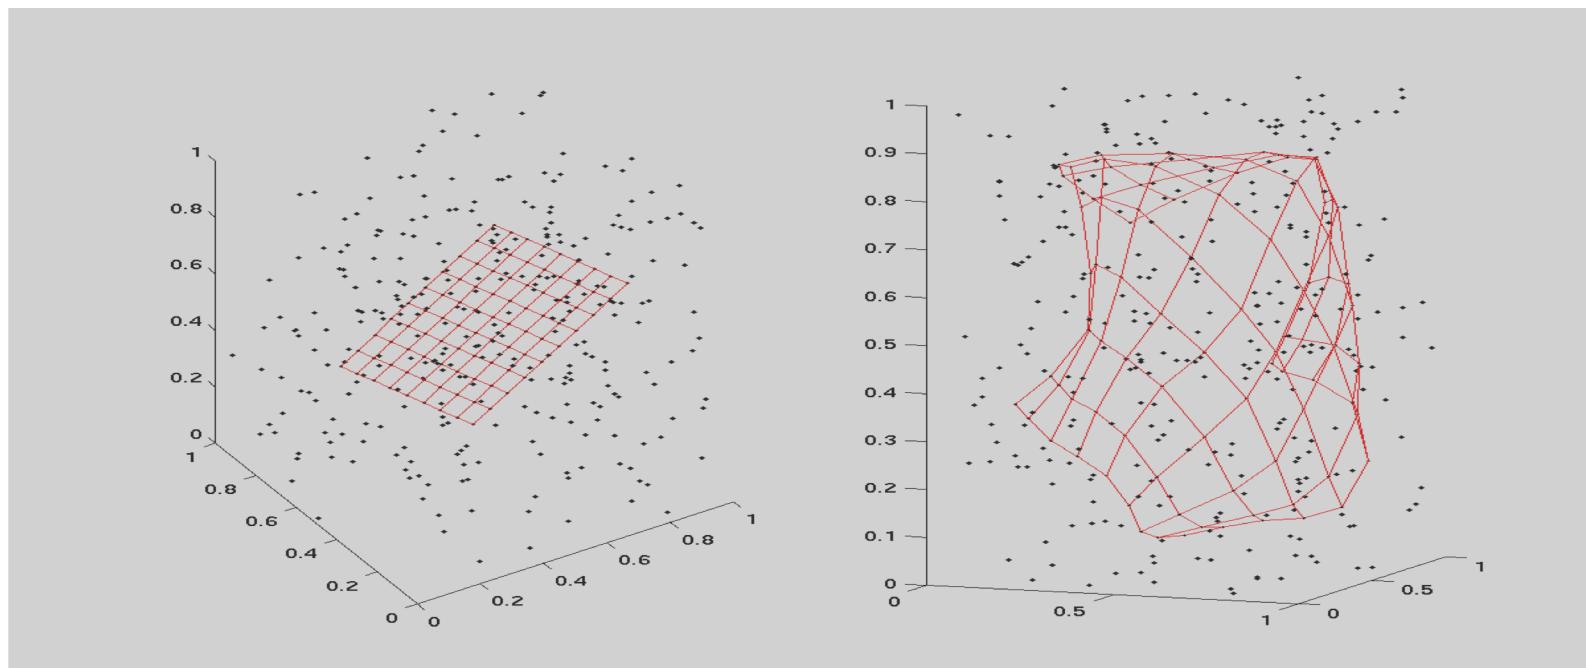
This is a divide & conquer approach



- **Projection**
 - Dimensionality compression
 - **Multi-dimensional Scaling**
 - Similitude information coding
- **Clustering**
 - Finding grouping structure in data
 - E.g., **Principal Components Analysis**
 - Similitude information coding
- ***Self-Organizing Map (SOM) & Generative Topographic Mapping (GTM)***
 - They combine latent representation and clustering

Projection

- Representation in <4-D, so that the distance-neighborhood relations between multi-dimensional points are faithfully preserved
 - It is impossible to preserve information integrally
 - Some scale normalization is required
- Linear vs. non-linear projections



Projection methods

- **Methods based on inter-point distances, where:**

dx = distance in the original space

dy = distance in the projection space

h = neighborhood function

$$E = \sum (dx - dy)^2$$

MDS, PCA

$$E = \sum (dx - dy)^2 / dx$$

Sammon's projection

$$E = \sum (dx - dy)^2 e^{-dy}$$

CCA

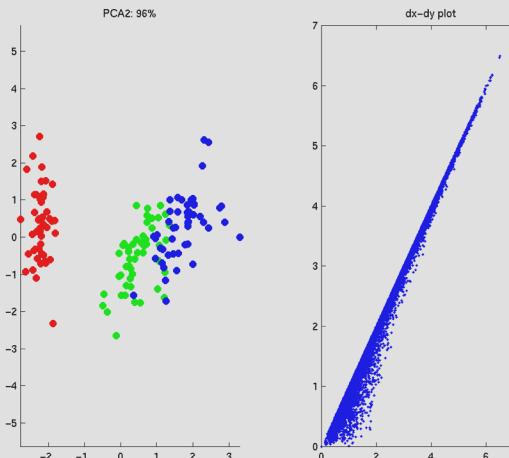
$$E = \sum dx^2 h(dy)$$

SOM

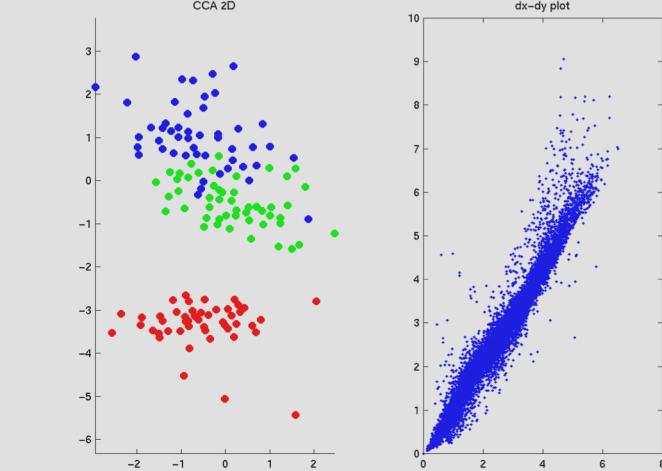
... and in which we aim to minimize an inherent projection distortion (E)

Example of Projection

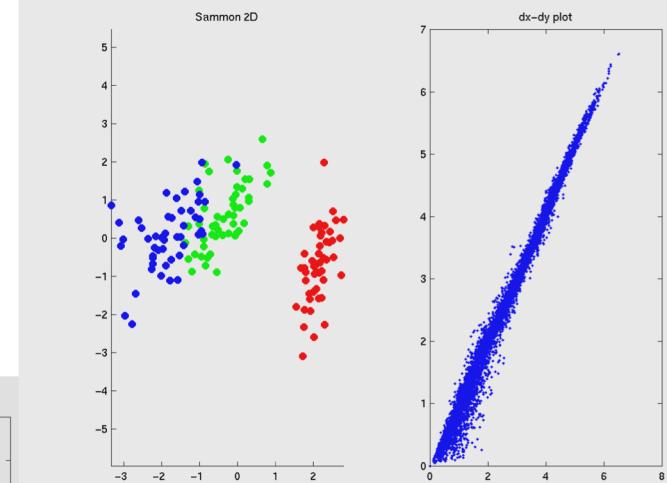
PCA



CCA



Sammon's projection





Self-Organizing Maps (SOM)



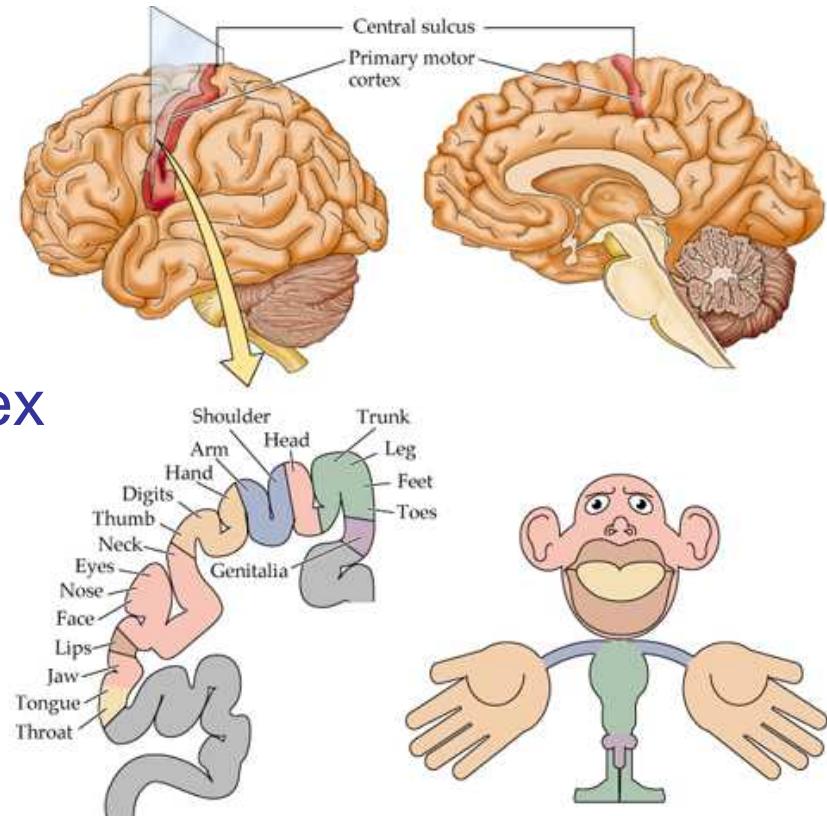
Look at the videos

<https://youtu.be/K4WuE7zIOZo>

https://youtu.be/0qtvb_Nx2tA

Introduction to Self-Organizing Maps

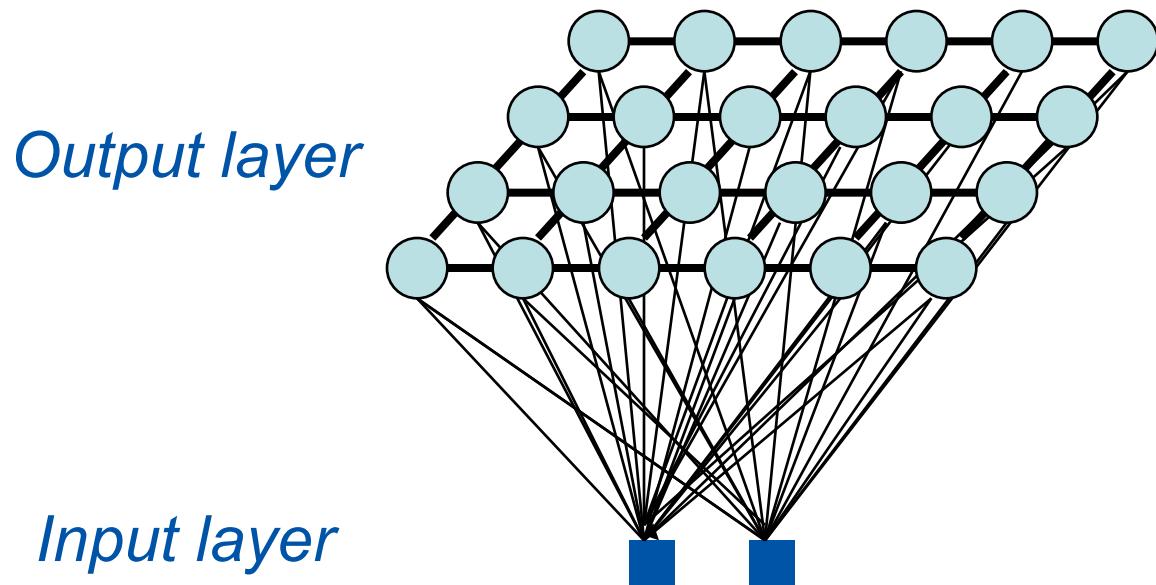
- T. Kohonen proposes SOM in 1982
- **SOM are simplified models of cortical maps in the brain**
- Things that are near in the outside world link to areas near in the cortex
- These areas are represented by two-dimensional maps
- Neighboring areas in these maps represent neighboring areas in the sensory input space



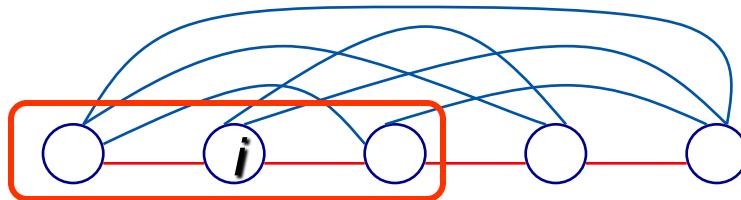
- Kohonen describes SOM as “**visualization and analysis tool for high dimensional data**”
- It is an unsupervised ANN
- Belongs to the category of competitive learning networks
- Applications
 - Clustering
 - Dimensionality reduction
 - Classification
 - Feature extraction (Self-Organized Feature Map)
 - Visualization

Self organizing maps

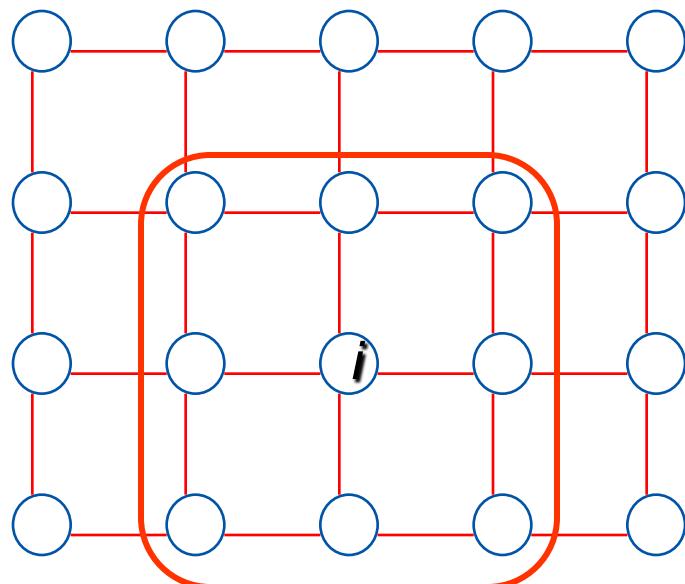
- The purpose of SOM is to map a multi-dimensional input space onto a **topology preserving** map of neurons
 - Preserve a topological so that neighboring neurons respond to «similar» input patterns
 - The topological structure is often a 2 or 3 dimensional space



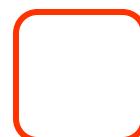
Commonly output-layer structures



One-dimensional
(completely interconnected
for determining “winner” unit)



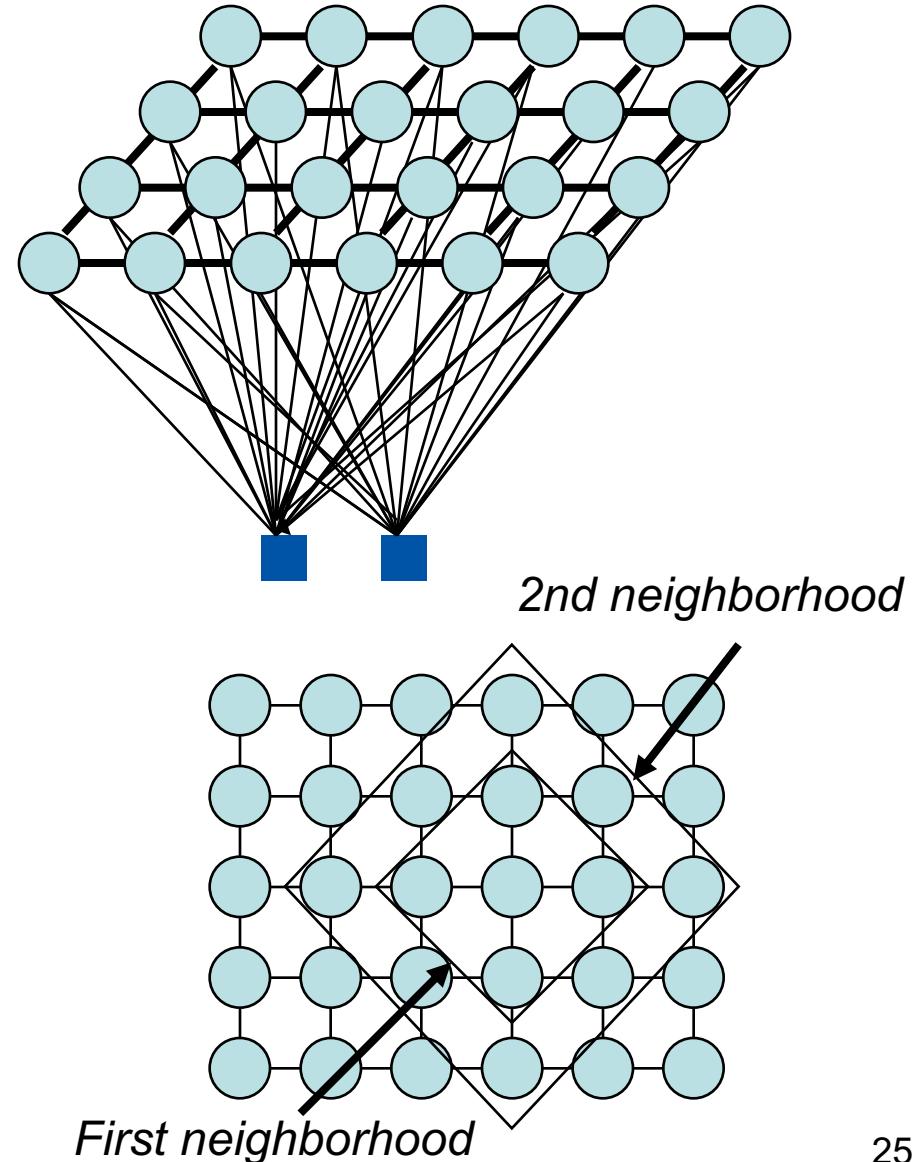
Two-dimensional
(connections omitted, only
neighborhood relations
shown [green])



Neighborhood of neuron i

Topological structure

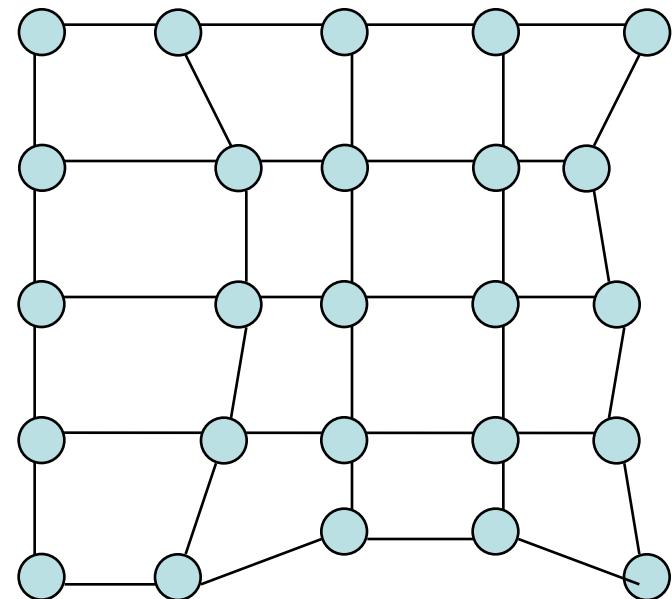
- The activation of the neuron is spread in its **direct neighborhood**
 - neighbors become sensitive to the same input patterns
- Block distance
- The size of the neighborhood is initially large but reduce over time
 - Specialization of the network



- SOM maps a multi-dimensional input space onto a topology preserving **map of neurons**
- Each neuron is assigned a weight vector with the same dimensionality of the input space
- Each Input pattern is compared to each weight vector
 - Distance is calculated between the input pattern and each neuron in the network (e.g. *Euclidean Distance*)
- Selects the neuron that is closest as the winning neuron

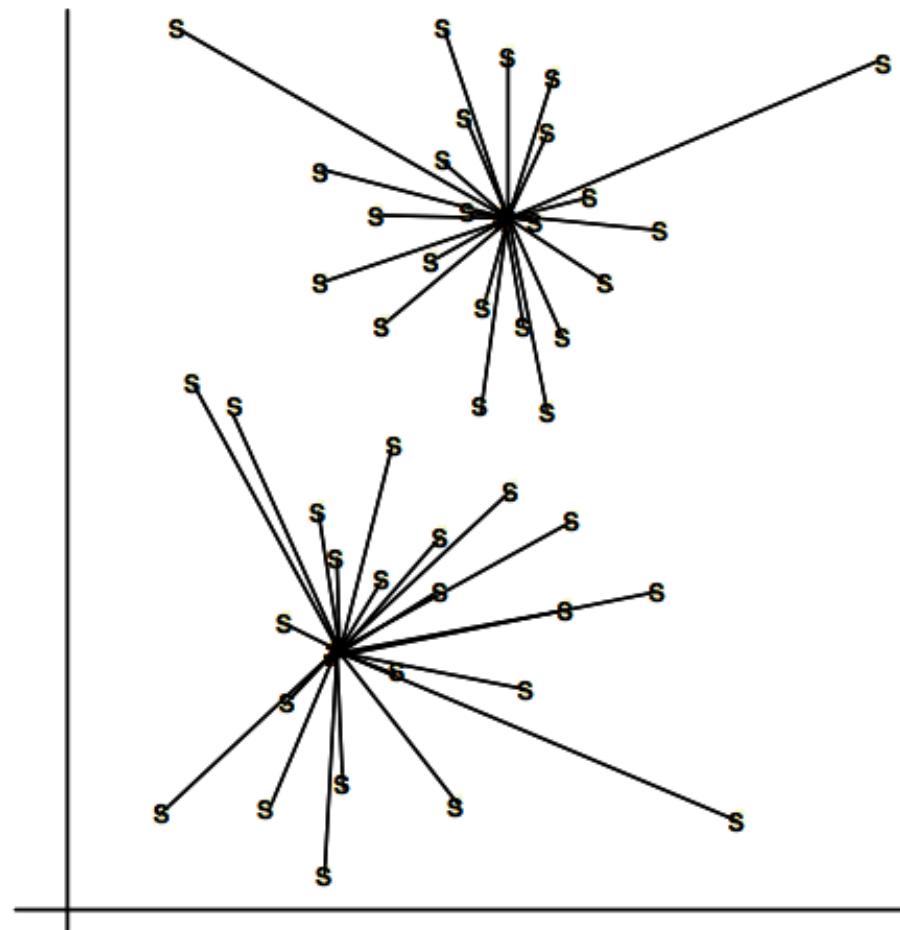
Adaptation in SOM

- During **training**, the “winner” neuron and its neighborhood adapts to make their weight vector more similar to the input pattern that caused the activation
- The neurons are moved closer to the input pattern
- The magnitude of the adaptation is controlled via a **learning parameter** which decays over time



Quantization error

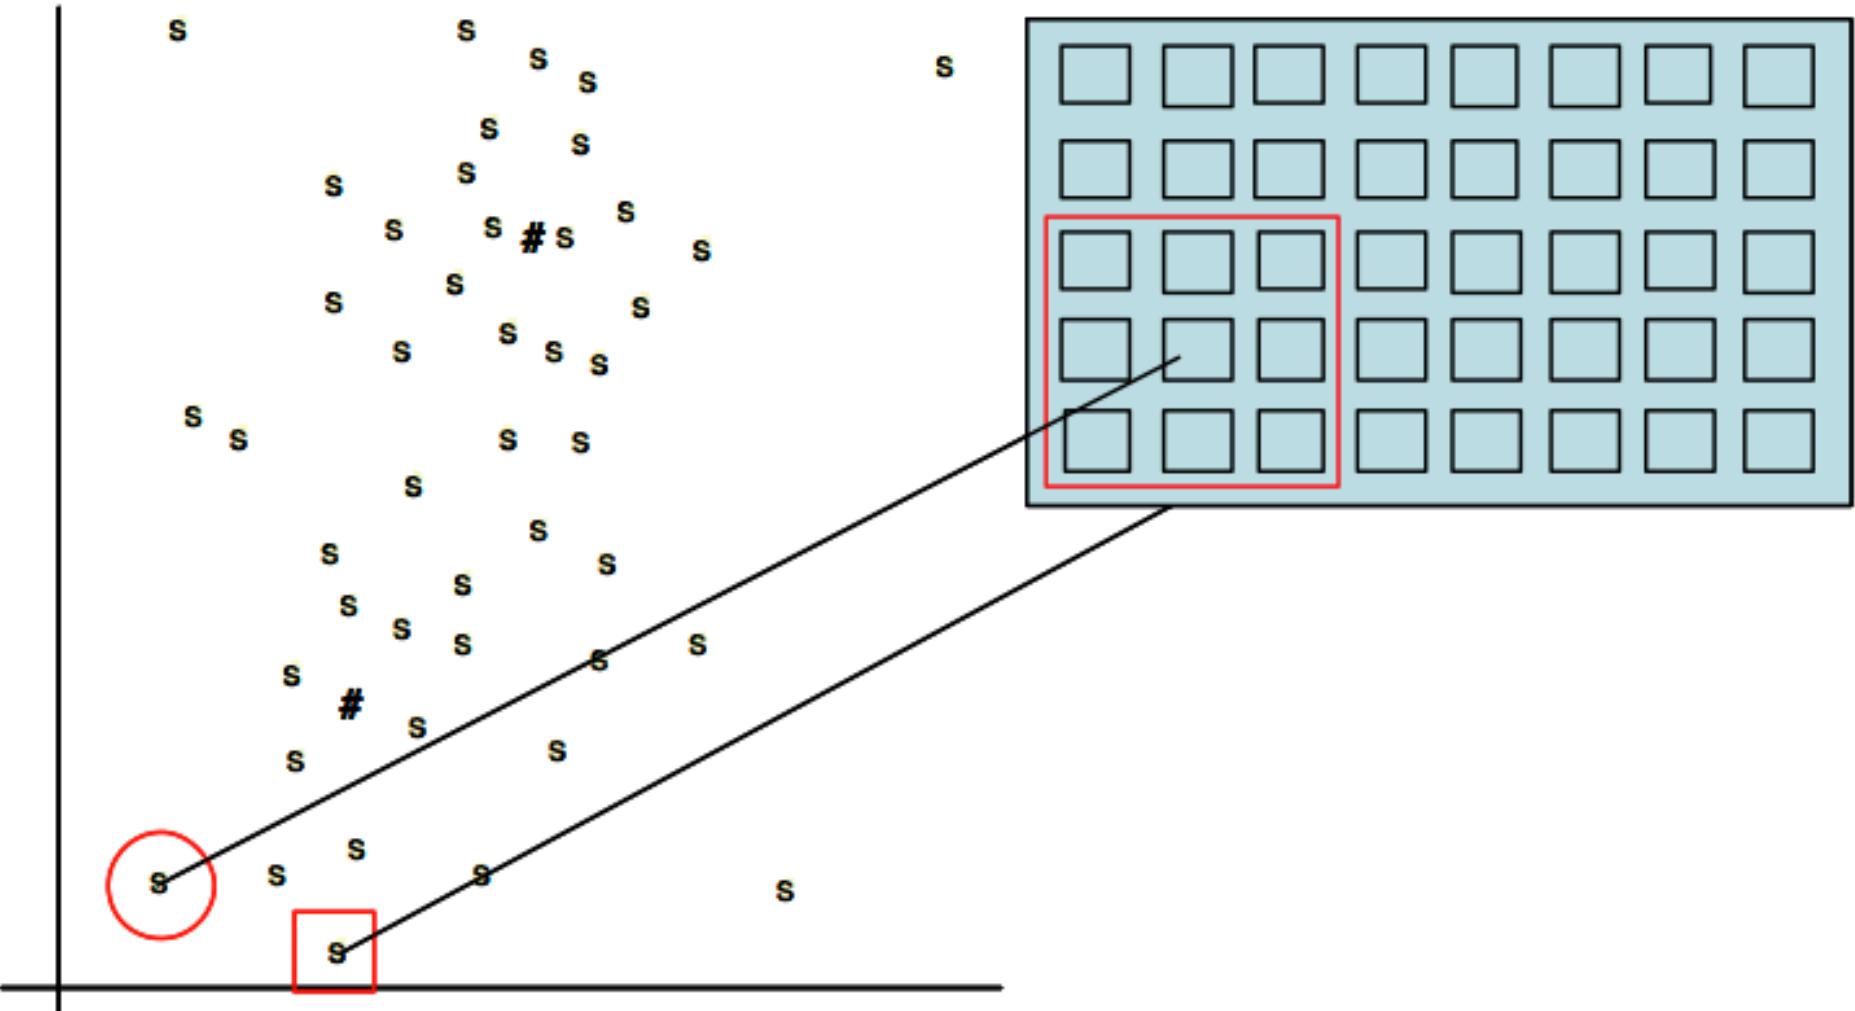
- Measures how well our neurons represent the input patterns
 - Note that there will be always some difference between the input pattern and the neuron it is mapped to
- It is calculated by summing all the distances between each input pattern and the neuron to which is mapped.



Topological error

- Evaluates the complexity of the output space
- Measures the number of times the second closest neighbor in the input space is not mapped into the neighbourhood of the neuron in the output space
- A high topological error may indicate that the classification problem is complex or may suggest that the training was not adequate and the network is folded

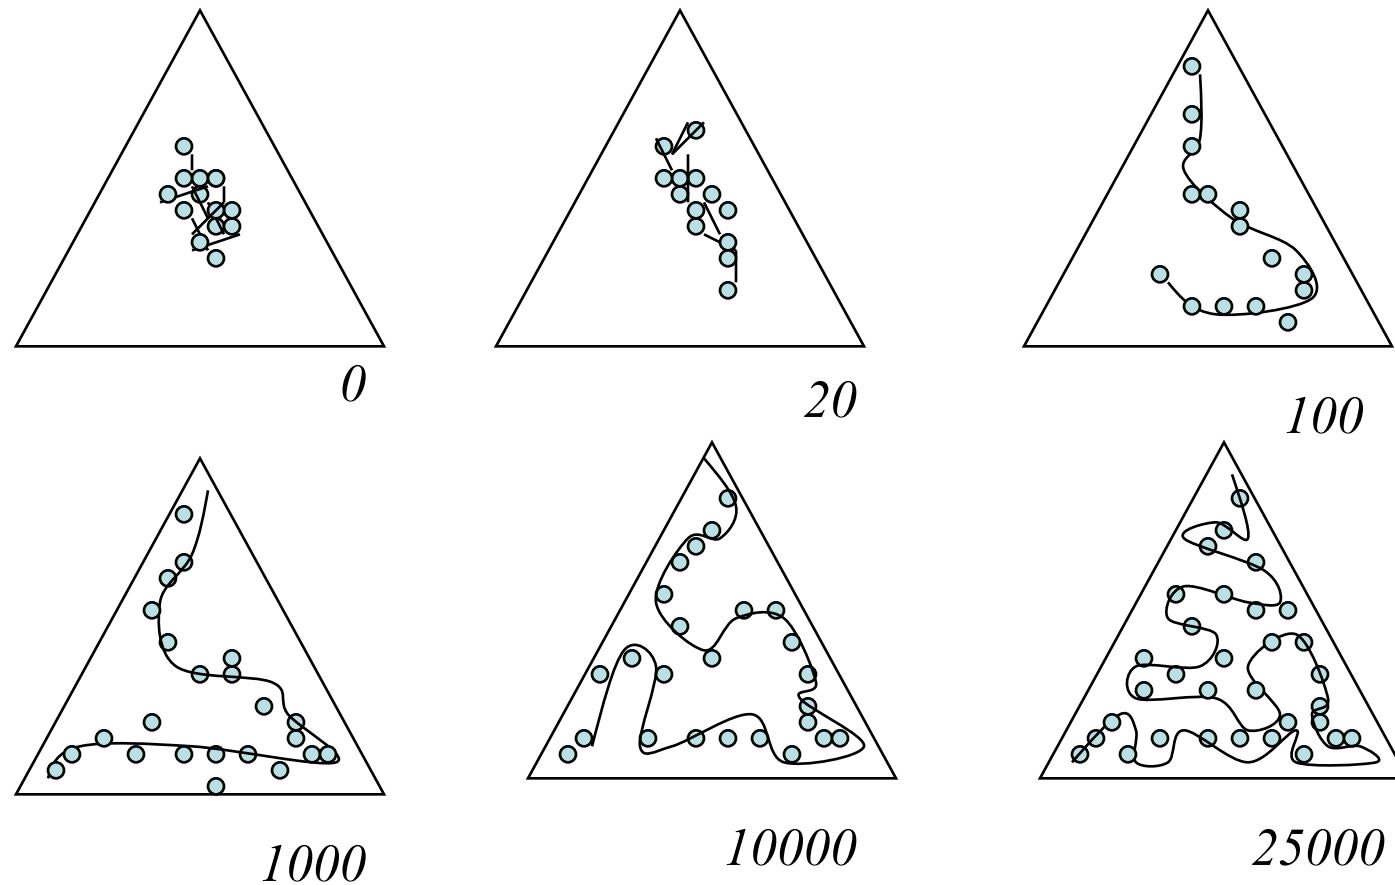
Topological error



- Before training, neurons may be initialized randomly
- **Part 1. Unfolding phase**
 - Neurons are “spread out” and pulled towards the general area (in the input space) where they will stay
- **Part 2. Fine tuning phase**
 - SOM match the neurons as far as possible to the input patterns, thus decreasing the quantization error

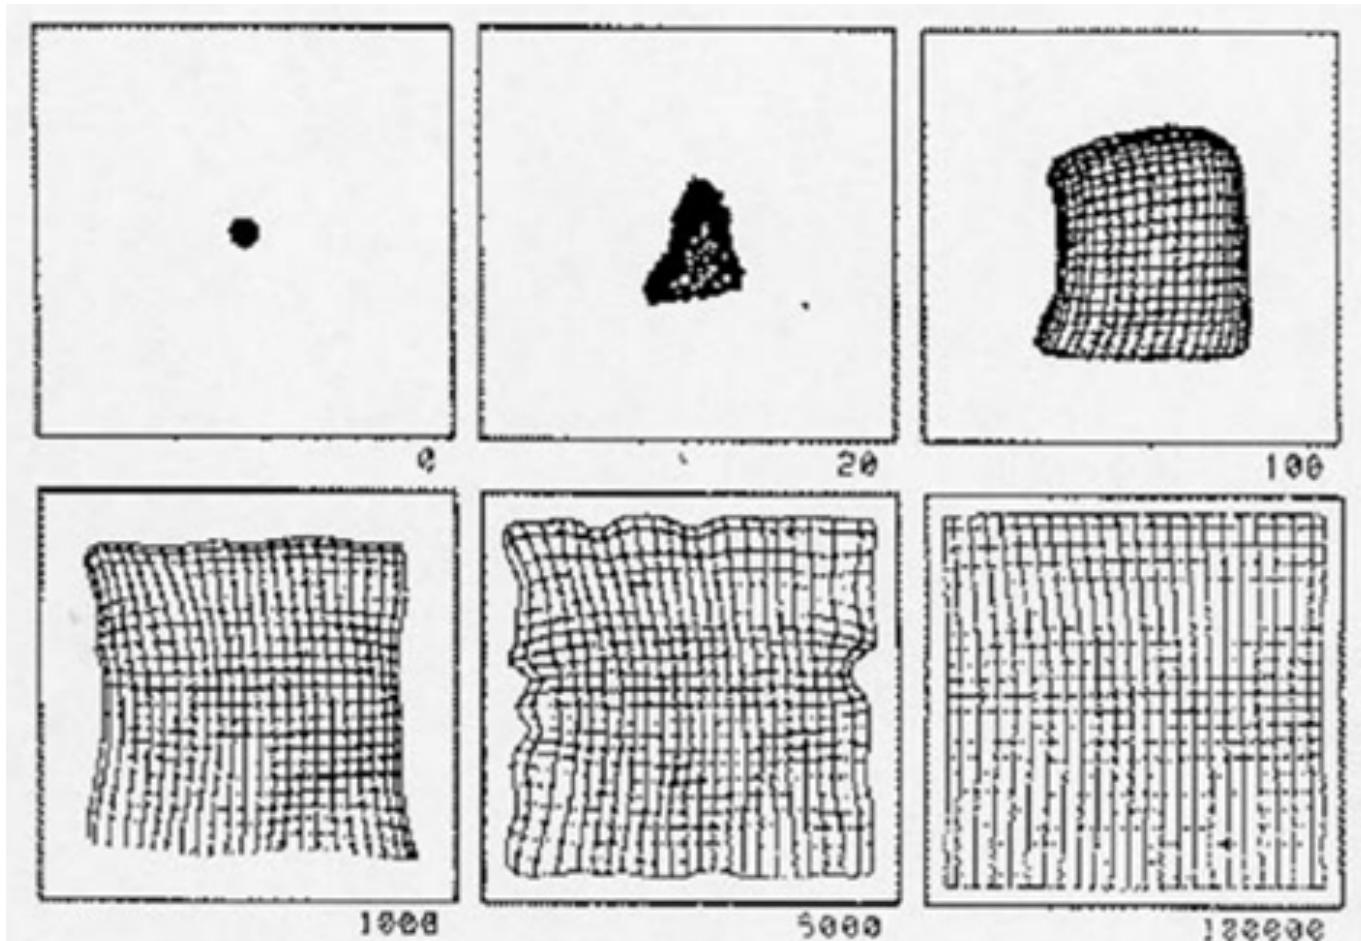
Example 1 when training SOM

Learning a one-dimensional representation of a two-dimensional (triangular) input space:



Example 2 when training a SOM

Learning a two-dimensional representation of a two-dimensional (square) input space:



SOM algorithm

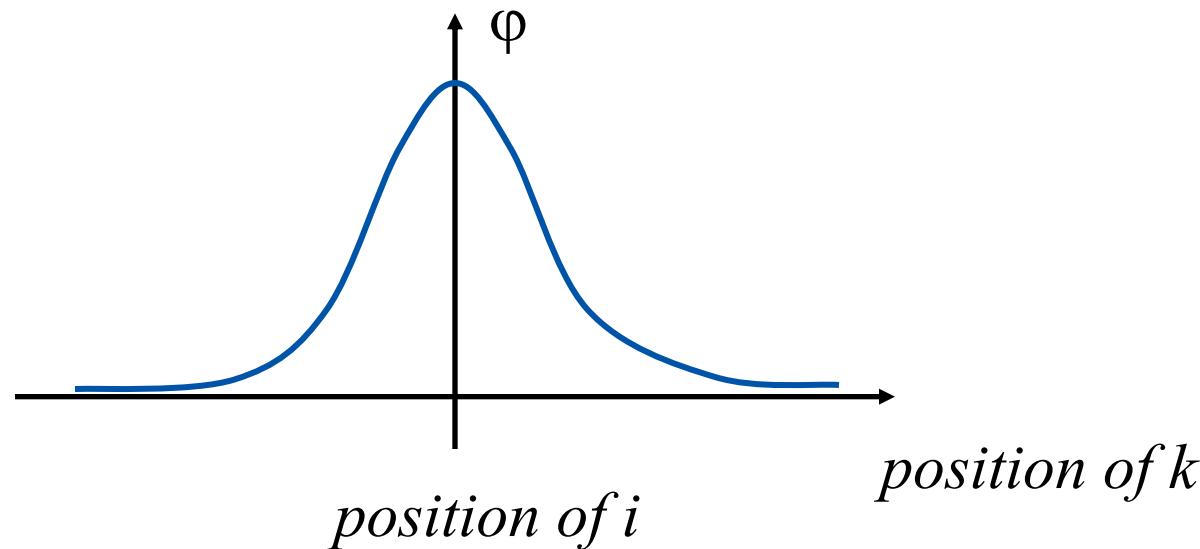
For n -dimensional input space and m output neurons:

1. Choose random weight vector w_i for neuron i , $i = 1, \dots, m$
2. For each input, choose random input x and do 2, 3, 4, and 5
3. Calculate the distance between the pattern and all neurons (usually the Euclidean distance)
4. Determine winner neuron k :
$$\|w_k - x\| = \min_i \|w_i - x\|$$
5. Update all weight vectors of all neurons i in the neighborhood of neuron k : $w_i := w_i + \eta \cdot \varphi(i, k) \cdot (x - w_i)$
(w_i is shifted towards x)
6. If convergence criterion met, STOP.
Otherwise, narrow neighborhood function φ and learning parameter η and go to (2).

Neighbourhood function

A neighborhood function $\varphi(i, k)$ indicates how closely neurons i and k in the output layer are connected to each other.

Usually, a Gaussian function on the distance between the two neurons in the layer is used:

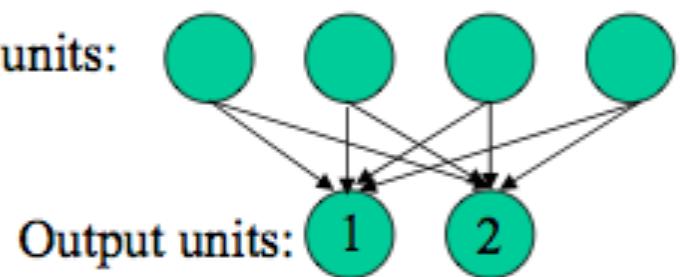


Example SOM

- From Fausett (1994)
- $n = 4, m = 2$
- Training samples
 - i1: (1, 1, 0, 0)
 - i2: (0, 0, 0, 1)
 - i3: (1, 0, 0, 0)
 - i4: (0, 0, 1, 1)

Network Architecture

Input units:



Output units:

What should we expect as outputs?

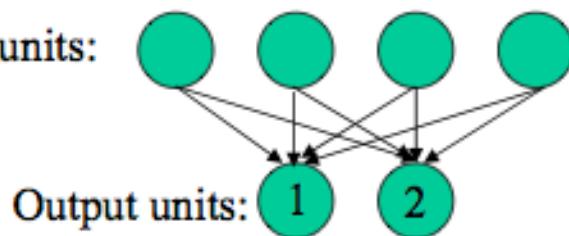
Distances between data samples

- Use Minkowski distance with $r = 1$
- Training samples

- i1: (1, 1, 0, 0)
- i2: (0, 0, 0, 1)
- i3: (1, 0, 0, 0)
- i4: (0, 0, 1, 1)

Network Architecture

Input units:



	i1	i2	i3	i4
i1	0			
i2	3	0		
i3	1	2	0	
i4	4	1	3	0

What might we expect from the SOM?

Example details

- Let Neighbourhood, $\varphi(i, k) = 0$
- Learning rate η
 - $\eta(t) = 0.6; 1 \leq t \leq 4$
 - $\eta(t) = 0.5 \eta(1); 5 \leq t \leq 8$
 - $\eta(t) = 0.5 \eta(5); 9 \leq t \leq 12$
- Initial weight matrix
 - Neuron 1, N1 [.2, .6, .5, .9]
 - Neuron 2, N2 [.8, .4, .7, .3]
- Weight update: $w_i(t+1) := w_i(t) + \eta(t) \cdot (x - w_i)$
- **Problem:** Calculate the weights updates for the first four steps

First weight update

- Training sample i1
- Distances (Euclidean) input pattern to all neurons
 - $D(1) = (.2-1)^2 + (.6-1)^2 + (.5-0)^2 + (.9-0)^2 = 1.86$
 - $D(2) = (.8-1)^2 + (.4-1)^2 + (.7-0)^2 + (.3-0)^2 = .98$
 - Neuron 2 (N2) wins
- Weights on winning neuron are updated
 - $w_i(t+1) := w_i(t) + .6 \cdot (x - w_i) = .4w_i(t) + .6x$
 - N1 [.2, .6, .5, .9]
 - N2 [.92, .76, .28, .12]

Second weight update

- Training sample i2
- Distances input pattern to all neurons
 - $D(1) = (.2-0)^2 + (.6-0)^2 + (.5-0)^2 + (.9-1)^2 = .66$
 - $D(2) = (.92-0)^2 + (.76-0)^2 + (.28-0)^2 + (.12-1)^2 = 2.28$
 - N1 wins
- Weights on winning neuron are updated
 - $w_i(t+1) := w_i(t) + .6 \cdot (x - w_i) = .4w_i(t) + .6x$
 - N1 [.08, .24, .20, .96]
 - N2 [.92, .76, .28, .12]

TODO: Complete the example with training i3 and i4

At the end of the example

- Once finished the first iteration (also known as epoch) the learning rate is updated, it will be .3
- Next, second iteration will start
- The process finishes when the number of iterations is reached
- After many iterations through the data set
 - N1 [0, 0, .5, 1]
 - N2 [1, .5, 0, 0]

Did we get the clustering we expected?

Solution

- Training samples

- i1: (1, 1, 0, 0)
- i2: (0, 0, 0, 1)
- i3: (1, 0, 0, 0)
- i4: (0, 0, 1, 1)

Sample: i1

- Weights

- N1 [0, 0, .5, 1]
- N2 [1, .5, 0, 0]

- Distance from unit1 weights

$$(1-0)^2 + (1-0)^2 + (0-.5)^2 + (0-1.0)^2 = 1+1+.25+1=3.25$$

- Distance from unit2 weights

$$(1-1)^2 + (1-.5)^2 + (0-0)^2 + (0-0)^2 = 0+.25+0+0=.25 \text{ (winner)}$$

Sample: i2

- Distance from unit1 weights

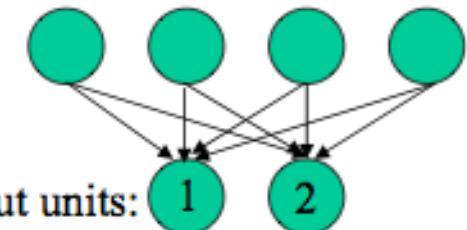
$$(0-0)^2 + (0-0)^2 + (0-.5)^2 + (1-1.0)^2 = 0+0+.25+0 \text{ (winner)}$$

- Distance from unit2 weights

$$(0-1)^2 + (0-.5)^2 + (0-0)^2 + (1-0)^2 = 1+.25+0+1=2.25$$

Network Architecture

Input units:



Output units: 1 2

Solution

- Training samples

- i1: (1, 1, 0, 0)
- i2: (0, 0, 0, 1)
- i3: (1, 0, 0, 0)
- i4: (0, 0, 1, 1)

Sample: i3

- Weights

- N1 [0, 0, .5, 1]
- N2 [1, .5, 0, 0]

- Distance from unit1 weights

$$(1-0)^2 + (0-0)^2 + (0-.5)^2 + (0-1.0)^2 = 1+0+.25+1=2.25$$

- Distance from unit2 weights

$$(1-1)^2 + (0-.5)^2 + (0-0)^2 + (0-0)^2 = 0+.25+0+0=.25 \text{ (winner)}$$

Sample: i4

- Distance from unit1 weights

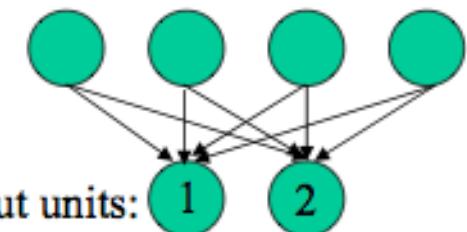
$$(0-0)^2 + (0-0)^2 + (1-.5)^2 + (1-1.0)^2 = 0+0+.25+0 \text{ (winner)}$$

- Distance from unit2 weights

$$(0-1)^2 + (0-.5)^2 + (1-0)^2 + (1-0)^2 = 1+.25+1+1=3.25$$

Network Architecture

Input units:



Output units: 1 2

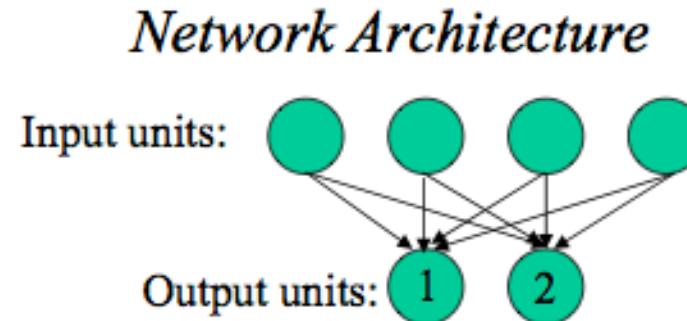
At the end of the example

- Training samples

- i1: (1, 1, 0, 0)
- i2: (0, 0, 0, 1)
- i3: (1, 0, 0, 0)
- i4: (0, 0, 1, 1)

- Weights

- N1 [0, 0, .5, 1]
- N2 [1, .5, 0, 0]



- Samples i1, i3 clusters with neuron 2
- Samples i2, i4 clusters with neuron 1

- **Also known as Kohonen Maps**
 - It is an **unsupervised ANN**
 - Cooperative learning (through neighbourhood function)
 - Competitive learning (*winner takes most –if not all-*)
 - Discretization (in the form of network grids) and projection are simultaneously performed
 - SOM provides a topology preserving mapping from the high dimensional space to map units
 - A SOM has the capability to generalize
 - ***k-means*** is an special case of **SOM**

What else?

Arriving at this point:



TO READ

An introduction to Self-Organizing Maps

- Read from page 299 to 315



In general all the chapter is interesting.

- You do not need to memorize the formulas, just to understand the concepts and what SOM is doing in its algorithm (the same contents that I have introduced you in the slides)

Multi-Dimensional Scaling (MDS)

Introduction

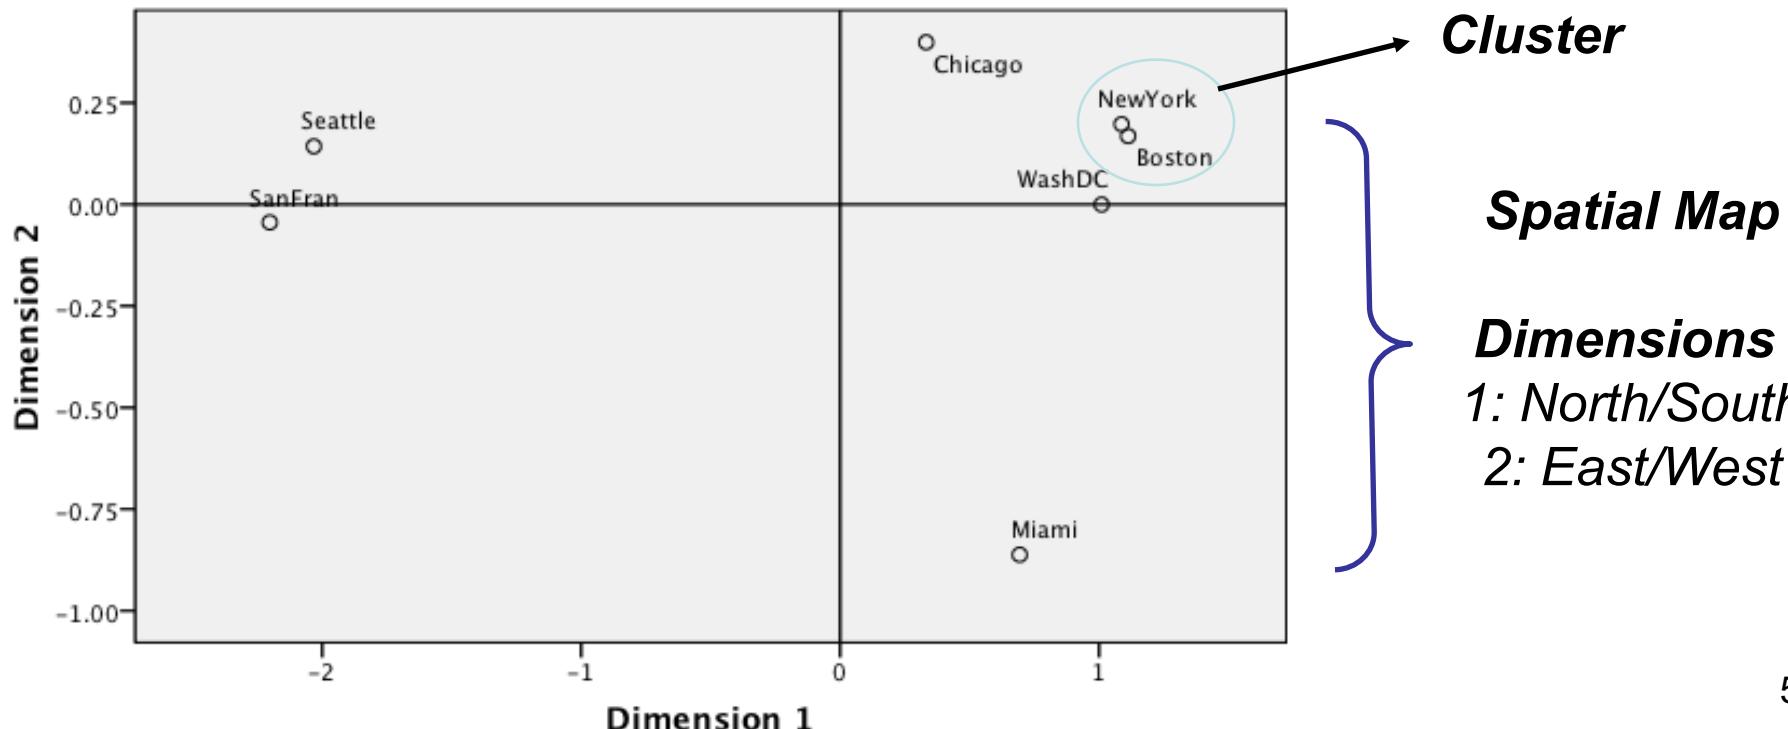
- “*If you are interested in how certain objects relate to each other ... and if you would like to present these relationships in the form of a map then MDS is the technique you need*”
- **The goal** of an MDS analysis is to find a spatial configuration of objects when all that is known is some measure of their general (dis)similarity.

Multidimensional Scaling

- Generally regarded as **exploratory data analysis**
- **Reduces large amounts of data** into easy-to-visualize structures
- Attempts to **find structure** (visual representation) in a set of distance measures, e.g. dis/similarities, between objects/instances
 - Shows how variables/objects are related perceptually
- **How?** By assigning instances to specific locations in space
- Distances between points in space match dis/similarities as closely as possible:
 - Similar objects: Close points
 - Dissimilar objects: Far apart points

**Distances
Matrix:
Symmetric**

		1	2	3	4	5	6	7	8	9
		BOST	NY	DC	MIAM	CHIC	SEAT	SF	LA	DENV
1	BOSTON	0	206	429	1504	963	2976	3095	2979	1949
2	NY	206	0	233	1308	802	2815	2934	2786	1771
3	DC	429	233	0	1075	671	2684	2799	2631	1616
4	MIAMI	1504	1308	1075	0	1329	3273	3053	2687	2037
5	CHICAGO	963	802	671	1329	0	2013	2142	2054	996
6	SEATTLE	2976	2815	2684	3273	2013	0	808	1131	1307
7	SF	3095	2934	2799	3053	2142	808	0	379	1235
8	LA	2979	2786	2631	2687	2054	1131	379	0	1059
9	DENVER	1949	1771	1616	2037	996	1307	1235	1059	0



The Process of MDS: The Data

- Data of MDS: **similarities, dissimilarities, distances, or proximities** reflects amount of dis/similarity or distance between pairs of objects.
- Distinction between similarity and dissimilarity data dependent on type of scale used:
 - **Dissimilarity** scale: Low # = high similarity & High # = high dissimilarity.
 - **Similarity** scale: Opposite of dissimilarity.
E.g. On a scale of 1-9 (1 being the same and 9 completely different) how similar are chocolate bars A and B? Dissimilarity scale.

MDS model classified according to:

1) Type of proximities:

- **Metric/quantitative:** Quantitative information/interval data about objects' proximities e.g. city distance
- **Non-metric/qualitative:** Qualitative information/nominal data about proximities e.g. rank order

2) Number of proximity matrices (distance, dis/similarity matrix).

- Proximity matrix is the input for MDS
- The above criteria yield:
 - 1) **Classical MDS:** One proximity matrix (metric or non-metric)
 - 2) **Replicated MDS:** Several matrices
 - 3) **Weighted MDS/Individual Difference Scaling:** Aggregate proximities and individual differences in a common MDS space

This course we will concentrate on the classical MDS. The next slides and the example come from this reading



TO READ

An introduction to MDS

- From page 4 to page 12 (Introduction, Section 1, Introduction Section 2, and Section 2.1)

The remaining of the article is OPTIONAL.

The MDS Model

- **Classical MDS** uses Euclidean principles to model data proximities in geometrical space, where distance (d_{ij}) between points i and j is defined as:

$$d_{ij} = \sqrt{\sum (x_{ia} - x_{ja})^2}$$

x_i and x_j specify coordinates of points i and j on dimension a, respectively.

- The modeled Euclidean distances are related to the observed proximities, δ_{ij} , by some transformation/function (f).
- Most MDS models assume that the data have the form:
$$\delta_{ij} = f(d_{ij}) = \sqrt{\sum (x_{ia} - x_{ja})^2}$$
- All MDS algorithms are a variation of the above

Classical MDS algorithm

We assume known a coordinate matrix X

1. Set up the matrix of squared proximities, $P^{(2)} = [p^2]$
2. Apply the double centering: $B = -\frac{1}{2}JP^{(2)}J'$
using the matrix $J = I - n^{-1}\mathbf{1}\mathbf{1}'$
where n is the number of objects.
3. Extract the m largest positive eigenvalues of B and the corresponding m eigenvectors
4. A m -dimensional spatial configuration of the n objects is derived from the coordinate matrix $X = E_m \Lambda_m^{1/2}$
where E_m is the matrix of m eigenvectors and Λ_m is the diagonal matrix of m eigenvalues

Classical MDS example

- Proximity matrix
 - Distances between 4 cities

	cph	aar	ode	aal
cph	0	93	82	133
aar	93	0	52	60
ode	82	52	0	111
aal	133	60	111	0
- Step 1 – Matrix of squared proximities

$$\mathbf{P}^{(2)} = \begin{bmatrix} 0 & 8649 & 6724 & 17689 \\ 8649 & 0 & 2704 & 3600 \\ 6724 & 2704 & 0 & 12321 \\ 17689 & 3600 & 12321 & 0 \end{bmatrix}$$

Classical MDS example

- Step 2- Apply the double centering

$$\mathbf{J} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} - 0.25 \times \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.75 & -0.25 & -0.25 & -0.25 \\ -0.25 & 0.75 & -0.25 & -0.25 \\ -0.25 & -0.25 & 0.75 & -0.25 \\ -0.25 & -0.25 & -0.25 & 0.75 \end{bmatrix}.$$

Applying \mathbf{J} to $\mathbf{P}^{(2)}$ yields the double centered matrix \mathbf{B}

$$\mathbf{B} = -\frac{1}{2}\mathbf{JP}^{(2)}\mathbf{J} = \begin{bmatrix} 5035.0625 & -1553.0625 & 258.9375 & -3740.938 \\ -1553.0625 & 507.8125 & 5.3125 & 1039.938 \\ 258.9375 & 5.3125 & 2206.8125 & -2471.062 \\ -3740.9375 & 1039.9375 & -2471.0625 & 5172.062 \end{bmatrix}.$$

Classical MDS example

- Step 3- Extract the two largest eigenvalues and the corresponding eigenvectors

$$\lambda_1 = 9724.168, \lambda_2 = 3160.986$$

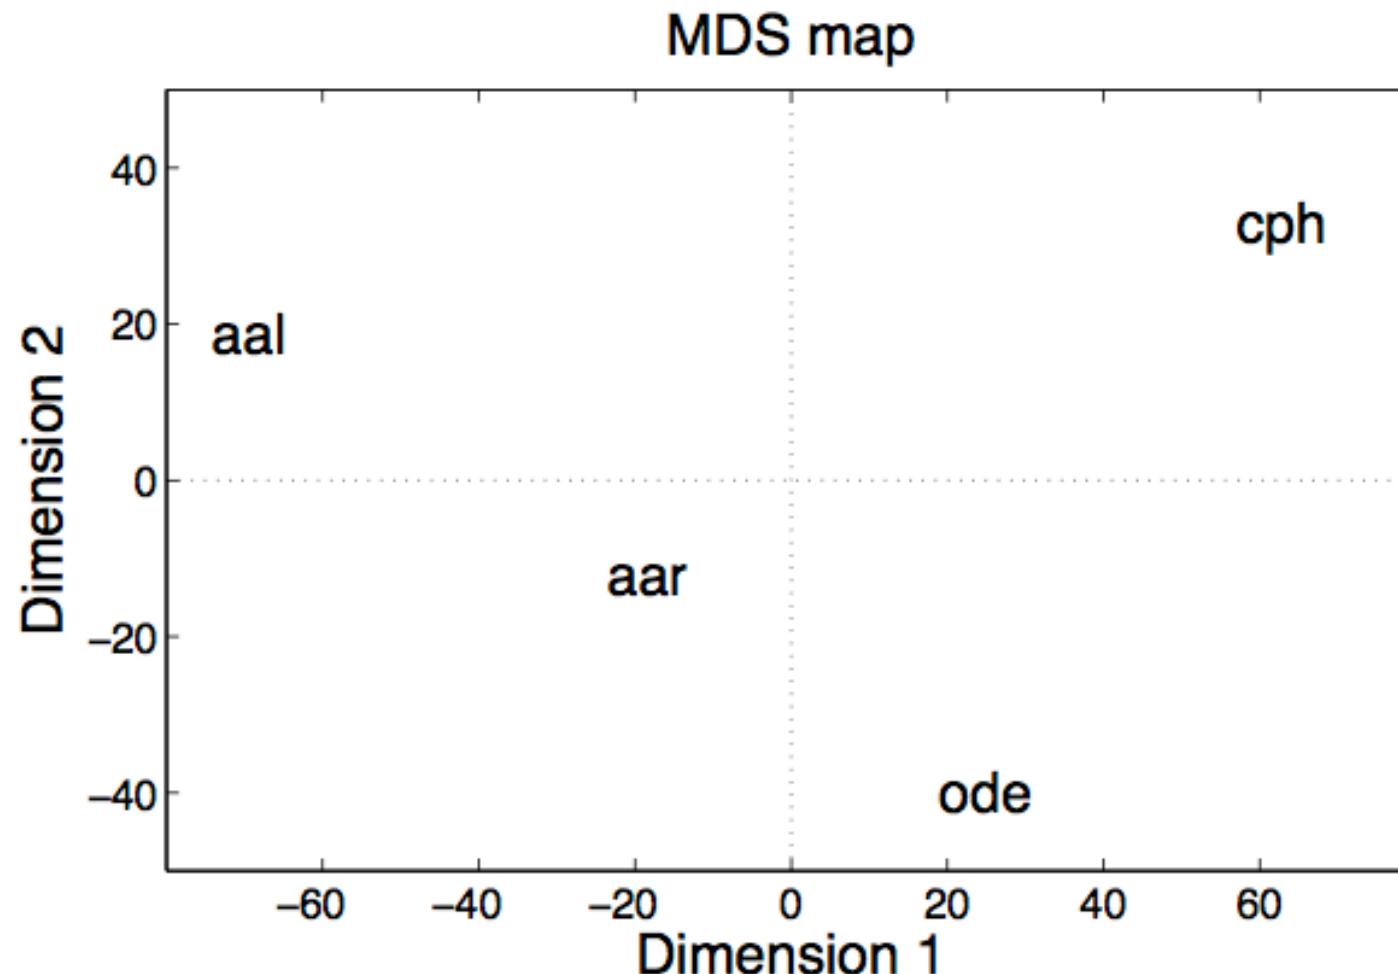
$$\mathbf{e}_1 = \begin{pmatrix} -0.637 \\ 0.187 \\ -0.253 \\ 0.704 \end{pmatrix}, \mathbf{e}_2 = \begin{pmatrix} -0.586 \\ 0.214 \\ 0.706 \\ -0.334 \end{pmatrix}$$

Classical MDS example

- Step 4- the coordinates of the cities are obtained by multiplying eigenvalues and vectors

$$\mathbf{X} = \begin{bmatrix} -0.637 & -0.586 \\ 0.187 & 0.214 \\ -0.253 & 0.706 \\ 0.704 & -0.334 \end{bmatrix} \begin{bmatrix} \sqrt{9724.168} & 0 \\ 0 & \sqrt{3160.986} \end{bmatrix} = \begin{bmatrix} -62.831 & -32.97448 \\ 18.403 & 12.02697 \\ -24.960 & 39.71091 \\ 69.388 & -18.76340 \end{bmatrix}$$

- The result is



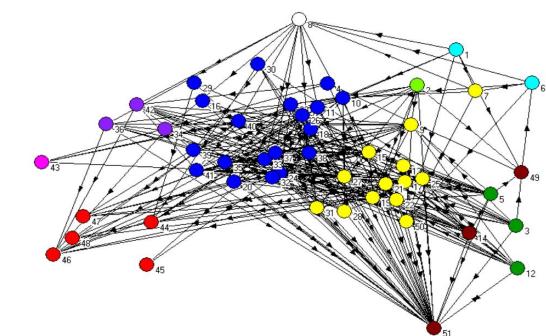
MDS Map/Perceptual Map/Spatial Representation:

1) **Clusters:** Groupings in a MDS spatial representation.

- These may represent a domain/subdomain.

2) **Dimensions:** Hidden structures in data. Ordered groupings that explain similarity between items.

- Axes are meaningless and orientation is arbitrary.
- In theory, there is no limit to the number of dimensions.
- In reality, the number of dimensions that can be perceived and interpreted is limited.



“Advantages” of MDS

- Does not require assumptions of linearity, metricity, or multivariate normality
- Can be used to model nonlinear relationships
- Dimensionality “solution” can be obtained from individuals; gives insight into how individuals differ from aggregate data
- Reveals dimensions without the need for defined attributes
- Dimensions that emerge from MDS can be incorporated into regression analysis to assess their relationship with other variables

“Disadvantages” of MDS

- Provides a global measure of dis/similarity but does not provide much insight into subtleties
- Increased dimensionality: Difficult to represent and decreases intuitive understanding of the data. As such, the model of the data becomes as complicated as the data itself.
- Determination of meanings of dimensions is subjective.

Summary

- **MDS:** technique used in data visualisation for exploring similarities or dissimilarities in data. An MDS algorithm starts with a matrix of item-item similarities, then assigns a location of each item in a low-dimensional space, suitable for graphing or 3D visualisation.
- **Taxonomy:**
 - Metric multidimensional scaling -- assumes the input matrix is just an item-item distance matrix. Analogous to PCA, an eigenvector problem is solved to find the locations that minimize distortions to the distance matrix. Its goal is to find a Euclidean distance approximating a given distance.
 - Generalized multidimensional scaling (GMDS) -- A superset of metric MDS that allows for the target distances to be non-Euclidean.
 - Non-metric multidimensional scaling -- It finds a non-parametric monotonic relationship between the dissimilarities in the item-item matrix and the Euclidean distance between items, and the location of each item in the low-dimensional space



EVERYBODY WHO WENT TO
THE MOON HAS EATEN
CHICKEN!



WHAT'S FREAKING US OUT HERE IS THAT WE'VE
FOUND A CORRELATION BETWEEN OWNING CATS
AND BEING STRUCK BY LIGHTNING.



Questions?

Week 5

Course. Introduction to Machine Learning

Theory 5. Visualization

Dr. Maria Salamó Llorente
maria.salamo@ub.edu

Dept. Mathematics and Informatics,
Faculty of Mathematics and Informatics,
University of Barcelona (UB)